

ESS 575: Occupancy Modeling Lab

Team England

09 November, 2022

Contents

Motivation	2
R libraries needed for this lab	2
Problem	3
Question 1	4
Question 2	4
Stochastic model of y (the observed detections):	4
Deterministic model of z (the true probability of occupancy):	4
Posterior and Joint:	5
Question 3	5
Posterior and Joint:	5
Question 4	6
Data set up	6
JAGS Model	6
Implement JAGS Model	7
Question 5	10
Question 6	11
Team England:	
• Caroline Blommel	
• Carolyn Coyle	
• Bryn Crosby	
• George Woolsey	

cblommel@mail.colostate.edu, carolynm@mail.colostate.edu, brcrosby@rams.colostate.edu, george.woolsey@colostate.edu

Motivation

Modeling presence or absence is a classic problem involving mixture models, specifically random variables that are zero-inflated. Extra zeros are encountered when we model presence or absence because zeros arise from two conditions: truly absent individuals and individuals present but undetected. This means we need a model for the process that controls occupancy, the true state, and model of the data that accounts for detection. This is often our starting point in Bayesian analysis – there is a true, unobserved state we seek to understand using a model of a process. We take imperfect observations on that state and must correct them using a model of the data.

R libraries needed for this lab

You need to load the following libraries. Set the seed to 10 to compare your answers to ours.

```
# bread-and-butter
library(tidyverse)
library(lubridate)
library(viridis)
library(scales)
library(latex2exp)
# visualization
library(cowplot)
library(kableExtra)
# jags and bayesian
library(rjags)
library(MCMCvis)
library(HDInterval)
library(BayesNSF)
#set seed
set.seed(10)
```

Problem

A fundamental question in landscape ecology seeks to understand how landscape structure shapes variation in habitat use by species. We will use data from the Swiss Survey of Common Breeding Birds, courtesy of Royle and Dorazio (2008), to model habitat occupancy by a resident bird in the Swiss Alps, the willow tit (*Parus montanus*). The data come from annual surveys of one km² quadrats distributed across Switzerland (Fig. 1). Surveys are conducted during the breeding season on three separate days, but some quadrats have missing data so that the number of replicate observations is fewer than three.

During each survey, an observer records every visual or acoustic detection of a breeding species (we do not differentiate between these two types of detection in this problem) and marks its location using a global positioning system or, in earlier years, a paper map. We assume that the true state (occupied or unoccupied) does not change among sample dates, an assumption known as closure. This assumption is reasonable because we are observing a resident species during the breeding season.




Fig. 1. The willow tit (left, credit: Francis C. Franklin) is one of 70 bird species that are surveyed annually for abundance in 267 1 km² sampling units distributed across Switzerland (right, credit: the Swiss Ornithological Institute).

We want to understand the influence of forest cover and elevation on the distribution of the willow tit. The data frame `BayesNSF::SwissBirds` has the number of times a quadrat (`quadrat`) was searched (`numberVisits`) and the number of times willow tits were detected (`numberDetections`). We have covariates on forest canopy cover (`forestCover`) as well as elevation in meters (`elevation`) for each quadrat surveyed. Data on detection each day on each quadrat (0 or 1) are also available. Develop a model of the influence of forest cover and elevation on the distribution of willow tits. Your model should allow estimation of the optimum elevation of willow tit habitat at the mean forest cover, where optimum elevation is defined as the elevation where probability of occupancy is maximum.

```
# data structure
BayesNSF::SwissBirds %>%
  dplyr::glimpse()

## Rows: 237
## Columns: 5
## $ quadrat      <int> 1, 2, 3, 4, 5, 6, 7, 8, 9, 10, 11, 12, 13, 14, 15, 16~
## $ numberVisits <int> 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, 3, ~
## $ numberDetections <int> 0, 0, 0, 0, 0, 0, 0, 0, 2, 1, 0, 0, 0, 0, 1, 2, 0, 3, ~
## $ elevation    <int> 420, 450, 1050, 1110, 510, 630, 590, 530, 1140, 770, ~
## $ forestCover  <int> 3, 21, 32, 35, 2, 60, 5, 13, 50, 57, 84, 17, 58, 26, ~
```

```
summary(BayesNSF::SwissBirds$quadrat)
```

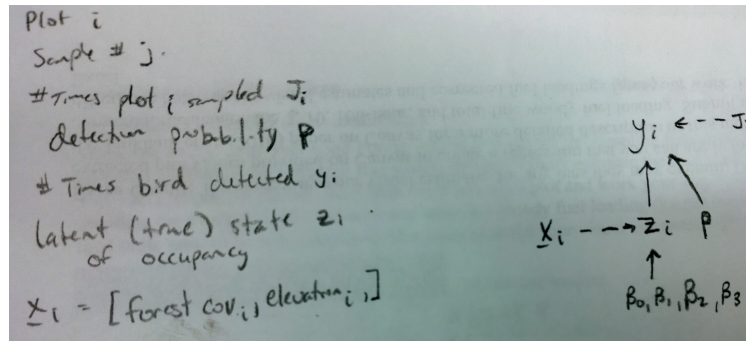
```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##         1       60     119     119    178     237
```

```
summary(BayesNSF::SwissBirds$numberVisits)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    1.000   3.000   3.000   2.806   3.000   3.000
```

Question 1

Diagram the network of knowns and unknowns.



Bayesian network for willow tit occupancy

Question 2

Write a mathematical expression for the posterior and the joint distribution.

Stochastic model of y (the observed detections):

$$y_i \sim \text{binomial}(p_i)$$

Deterministic model of z (the true probability of occupancy):

$$z_i \sim \text{Bernoulli}(g(\beta_0, \beta_1, \beta_2, \beta_3, x_i))$$

$$g(\beta_0, \beta_1, \beta_2, \beta_3, x_i) = \text{inverse logit}(\beta_0 + \beta_1 x_{1i} + \beta_2 x_{1i}^2 + \beta_3 x_{2i}) = \frac{\exp(\beta_0 + \beta_1 x_{1i} + \beta_2 x_{1i}^2 + \beta_3 x_{2i})}{1 + \exp(\beta_0 + \beta_1 x_{1i} + \beta_2 x_{1i}^2 + \beta_3 x_{2i})}$$

$x_1 = \text{elevation}; x_2 = \text{forest cover}$

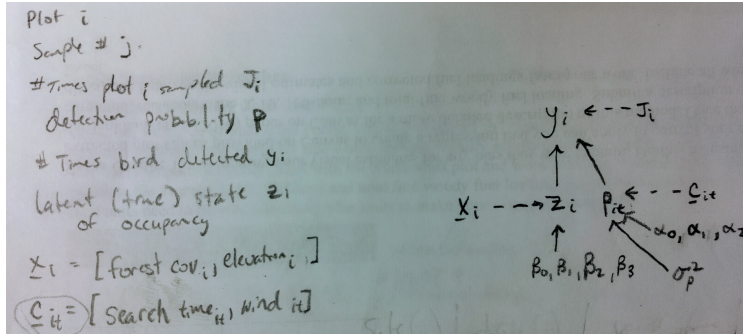
Utilizing a quadratic regression model with the inclusion of x_{1i}^2 in the model allows us to account for a nonlinear (e.g. parabolic) relationship between occupancy and elevation. This model form allows us to estimate the optimum elevation of willow tit habitat at the mean forest cover.

Posterior and Joint:

$$[z, \beta, p | \mathbf{y}] \propto \prod_{i=1}^{n=237} \text{binomial}(y_i | J_j, z_i \cdot p) \cdot \text{Bernoulli}(z_i | g(\beta, \mathbf{x}_i)) \cdot \text{uniform}(p | 0, 1) \\ \times \prod_{j=1}^{J=3} \text{normal}(\beta | 0, 2.7)$$

Question 3

Modify your model to include the effect of search time and wind speed (measured at each quadrat on each day) on detection probability. Draw a DAG and write the posterior and joint distributions. In so doing, assume that posterior predictive checks of a preliminary detection model revealed that you need to include an explicit variance term for the detection probability.



Bayesian network for willow tit occupancy with detection probability variance

Deterministic model of p (the probability of detection):

$$p_{it} \sim \text{beta}(\mu, \sigma^2) \\ \mu_{it} = g_2(\alpha_0, \alpha_1, \alpha_2, x_{it}) = \text{inverse logit}(\alpha_0 + \alpha_1 x_{1it} + \alpha_2 x_{2it}) = \frac{\exp(\alpha_0 + \alpha_1 x_{1it} + \alpha_2 x_{2it})}{1 + \exp(\alpha_0 + \alpha_1 x_{1it} + \alpha_2 x_{2it})} \\ x_1 = \text{search time}; x_2 = \text{wind speed}$$

Posterior and Joint:

$$[z, \beta, \alpha, p, \sigma^2 | \mathbf{y}] \propto \prod_{i=1}^{n=237} \prod_{t=1}^T \text{Bernoulli}(y_i | J_j, z_i \cdot p) \cdot \text{Bernoulli}(z_i | g(\beta, \mathbf{x}_i)) \cdot \text{beta}(p_{it} | g_2(\alpha, \mathbf{x}_{it}), \sigma^2) \\ \times \prod_{j=1}^{J=3} \text{normal}(\beta | 0, 2.7) \\ \times \prod_{j=1}^{J=3} \text{normal}(\alpha | 0, 2.7) \cdot \text{uniform}(\sigma^2 | 0, 100)$$

Question 4

Approximate the marginal posterior distributions of parameters in the forest and elevation model with constant detection probability (the first one, above) using JAGS. Conduct posterior predictive checks. Some hints: 1) You will need to standardize the covariates by subtracting the mean and dividing by the standard deviation for each observation in the elevation and forest cover data. Use the `scale` function to do this (it will drastically speed convergence). 2) You *must* give initial values of 1 to all unknown 0 or 1 z states.

Data set up

```
mean_elevation <- mean(BayesNSF::SwissBirds$elevation)
sd_elevation <- sd(BayesNSF::SwissBirds$elevation)
mean_forestCover <- mean(BayesNSF::SwissBirds$forestCover)
sd_forestCover <- sd(BayesNSF::SwissBirds$forestCover)

# list of data
data = list(
  N = nrow(BayesNSF::SwissBirds) # number sites
  , y = as.double(BayesNSF::SwissBirds$numberDetections)
  , elevation = scale(BayesNSF::SwissBirds$elevation) %>% as.numeric() %>% as.double
  , forestCover = scale(BayesNSF::SwissBirds$forestCover) %>% as.numeric() %>% as.double
  , n = as.double(BayesNSF::SwissBirds$numberVisits)
)
```

JAGS Model

Write out the JAGS code for the model.

```
## JAGS Model
model{
  # priors
  b0 ~ dnorm(0, (1/2.7)) # dnorm(mu = mean, tau= precision )
  b1 ~ dnorm(0, (1/2.7))
  b2 ~ dnorm(0, (1/2.7))
  b3 ~ dnorm(0, (1/2.7))
  p ~ dunif(0, 1)
  # likelihood
  for(i in 1:N){
    # deterministic model
    psi[i] <- ilogit(
      b0
      + b1*elevation[i]
      + b2*(elevation[i]^2)
      + b3*forestCover[i]
    )
    # Deterministic model of z (the true probability of occupancy)
    z[i] ~ dbern(psi[i])
    # likelihood
    # Stochastic model of y (the observed detections)
    y[i] ~ dbin(z[i] * p, n[i])
    y_sim[i] ~ dbin(z[i] * p, n[i])
  }
}
```

```

}
# Derived quantities
#posterior predictive checks
mean_y = mean(y)
sd_y = sd(y)
mean_y_sim = mean(y_sim)
sd_y_sim = sd(y_sim)
p_val_mean = step(mean_y_sim - mean_y)
p_val_sd = step(sd_y_sim - sd_y)
# sum of squares
for(j in 1:N) {
  sq[j] <- (y[j]-psi[j])^2
  sq_sim[j] <- (y_sim[j]-psi[j])^2
}
fit_y <- sum(sq)
fit_y_sim <- sum(sq_sim)
p_val_fit = step(fit_y_sim - fit_y)
}

```

Implement JAGS Model

```

#####
# insert JAGS model code into an R script
#####
{ # Extra bracket needed only for R markdown files - see answers
  sink("BirdOccupancyJAGS.R") # This is the file name for the jags code
  cat("
model{
  # priors
  b0 ~ dnorm(0, (1/2.7)) # dnorm(mu = mean, tau= precision )
  b1 ~ dnorm(0, (1/2.7))
  b2 ~ dnorm(0, (1/2.7))
  b3 ~ dnorm(0, (1/2.7))
  p ~ dunif(0, 1)
  # likelihood
  for(i in 1:N){
    # deterministic model
    psi[i] <- ilogit(
      b0
      + b1*elevation[i]
      + b2*(elevation[i]^2)
      + b3*forestCover[i]
    )
    # Deterministic model of z (the true probability of occupancy)
    z[i] ~ dbern(psi[i])
    # likelihood
    # Stochastic model of y (the observed detections)
    y[i] ~ dbin(z[i] * p, n[i])
    y_sim[i] ~ dbin(z[i] * p, n[i])
  }
  # Derived quantities
  #posterior predictive checks

```

```

    mean_y = mean(y)
    sd_y = sd(y)
    mean_y_sim = mean(y_sim)
    sd_y_sim = sd(y_sim)
    p_val_mean = step(mean_y_sim - mean_y)
    p_val_sd = step(sd_y_sim - sd_y)
    # sum of squares
    for(j in 1:N) {
      sq[j] <- (y[j]-psi[j])^2
      sq_sim[j] <- (y_sim[j]-psi[j])^2
    }
    fit_y <- sum(sq)
    fit_y_sim <- sum(sq_sim)
    p_val_fit = step(fit_y_sim - fit_y)
  }
  ", fill = TRUE)
sink()
}

#####
# implement model
#####
# specify the initial conditions for the MCMC chain
inits = list(
  list(
    b0 = runif(n = 1, min = -3, max = 3)
    , b1 = runif(n = 1, min = -3, max = 3)
    , b2 = runif(n = 1, min = -3, max = 3)
    , b3 = runif(n = 1, min = -3, max = 3)
    , p = runif(n = 1, min = 0, max = 1)
    , z = rep(1, nrow(BayesNSF::SwissBirds))
  )
  , list(
    b0 = runif(n = 1, min = -3, max = 3)
    , b1 = runif(n = 1, min = -3, max = 3)
    , b2 = runif(n = 1, min = -3, max = 3)
    , b3 = runif(n = 1, min = -3, max = 3)
    , p = runif(n = 1, min = 0, max = 1)
    , z = rep(1, nrow(BayesNSF::SwissBirds))
  )
  , list(
    b0 = runif(n = 1, min = -3, max = 3)
    , b1 = runif(n = 1, min = -3, max = 3)
    , b2 = runif(n = 1, min = -3, max = 3)
    , b3 = runif(n = 1, min = -3, max = 3)
    , p = runif(n = 1, min = 0, max = 1)
    , z = rep(1, nrow(BayesNSF::SwissBirds))
  )
)
# specify 3 scalars, n.adapt, n.update, and n.iter
# n.adapt = number of iterations that JAGS will use to choose the sampler
# and to assure optimum mixing of the MCMC chain
n.adapt = 1000
# n.update = number of iterations that will be discarded to allow the chain to

```



```

#   converge before iterations are stored (aka, burn-in)
n.update = 10000
# n.iter = number of iterations that will be stored in the
#   final chain as samples from the posterior distribution
n.iter = 10000
#####
# Call to JAGS
#####
jm = rjags::jags.model(
  file = "BirdOccupancyJAGS.R"
  , data = data
  , inits = inits
  , n.chains = length(inits)
  , n.adapt = n.adapt
)

```

```

## Compiling model graph
##   Resolving undeclared variables
##   Allocating nodes
## Graph information:
##   Observed stochastic nodes: 237
##   Unobserved stochastic nodes: 479
##   Total graph size: 3546
##
## Initializing model

```

```

stats::update(jm, n.iter = n.update, progress.bar = "none")
# save the coda object (more precisely, an mcmc.list object) to R as "zm"
zm = rjags::coda.samples(
  model = jm
  , variable.names = c(
    "b0"
    , "b1"
    , "b2"
    , "b3"
    , "mean_y"
    , "sd_y"
    , "mean_y_sim"
    , "sd_y_sim"
    , "p_val_mean"
    , "p_val_sd"
    , "fit_y"
    , "fit_y_sim"
    , "p_val_fit"
  )
  , n.iter = n.iter
  , n.thin = 1
  , progress.bar = "none"
)
#####
# check output
#####
# summary

```

```

MCMCvis::MCMCsummary(zm, params = c(
  "b0"
  , "b1"
  , "b2"
  , "b3"
  , "mean_y"
  , "mean_y_sim"
  , "p_val_mean"
  , "sd_y"
  , "sd_y_sim"
  , "p_val_sd"
  , "fit_y"
  , "fit_y_sim"
  , "p_val_fit"
)
)

```

```

##              mean          sd        2.5%        50%        97.5% Rhat
## b0          -0.2067444  0.27209676 -0.7402051 -0.2081981  0.3301068    1
## b1           1.9641654  0.30099666  1.4065907  1.9508201  2.5920371    1
## b2          -1.0890688  0.26684480 -1.6374844 -1.0798408 -0.5860213    1
## b3           0.8485297  0.23311605  0.4047594  0.8421950  1.3208771    1
## mean_y       0.7679325  0.00000000  0.7679325  0.7679325  0.7679325   NaN
## mean_y_sim   0.7658785  0.03708423  0.6877637  0.7679325  0.8354430    1
## p_val_mean   0.5131000  0.49983669  0.0000000  1.0000000  1.0000000    1
## sd_y         1.1900653  0.00000000  1.1900653  1.1900653  1.1900653   NaN
## sd_y_sim     1.1544843  0.04209644  1.0674191  1.1558458  1.2314852    1
## p_val_sd     0.2037667  0.40280420  0.0000000  0.0000000  1.0000000    1
## fit_y        291.9338309  8.70409156 276.4691540 291.4197765 310.3878655    1
## fit_y_sim    282.9456728 27.03912780 229.4597448 282.9275952 335.5894211    1
## p_val_fit     0.3699000  0.48278542  0.0000000  0.0000000  1.0000000    1
##              n.eff
## b0              4472
## b1              4623
## b2              3093
## b3              6787
## mean_y           0
## mean_y_sim    22025
## p_val_mean    24947
## sd_y           0
## sd_y_sim     19344
## p_val_sd     25934
## fit_y        11430
## fit_y_sim    17394
## p_val_fit    24742

```

Question 5

Summarize the parameters and check chains for convergence. Exclude the predictions of ψ from the summary if they were included in the coda object. What can you conclude about model fit?

```
# asdf
```

Question 6

What can you conclude about the relative importance of elevation and forest cover in controlling the bird's distribution? Plot the median probability of occupancy and the 95% highest posterior density interval as a function of elevation at the mean of forest cover. Find the optimum elevation of willow tit habitat at the mean forest cover, where optimum elevation is defined as the elevation where probability of occupancy is maximum. Plot a normalized histogram of MCMC output for the optimum elevation at the average forest cover. Overlay 0.95 highest posterior density limits on the optimum elevation.